ALGORITHMS & DATA STRUCTURES

CSCI 382

Erica Blum

2023

Notes by Aliya Ghassaei

Contents

1 Asymptotic notation

1.1 Definitions

1.1.1 Big-O

Intuition: If $f(n) \in O(g(n))$ then $f(n)$ grows no faster than $g(n)$. g is an asymptotic upper bound for f.

Definition 1.1. $O(g(n)) := \{f(n) : \exists\, c, n_0 > 0 \mid 0 \le f(n) \le c \cdot g(n) \text{ for all } n \ge n_0\}$

1.1.2 Big-Ω

Intuition: If $f(n) \in \Omega(g(n))$, $f(n)$ grows at least as fast as $g(n)$. f is a lower bound on f.

Definition 1.2. $\Omega(g(n)) := \{f(n) : \exists\, c, n_0 > 0 \mid 0 \le c \cdot g(n) \le f(n) \text{ for all } n \ge n_0\}$

1.1.3 Big-Θ

Intuition: $f(n)$ grows at the same rate/within a constant factor of $g(n)$

Definition 1.3. $\Theta(g(n)) := \{f(n) : \exists\, c_1, c_2, n_0 > 0 \mid 0 \le c_1 \cdot g(n) \le f(n) \le c_2 \cdot g(n) \text{ for all } n \ge n_0\}$

Theorem 1.1. For any two functions $f(n)$ and $g(n)$, $f(n) = \Theta(g(n)) \iff f(n) = O(g(n))$ and $f(n) = \Omega(g(n))$

1.1.4 Little-o

Definition 1.4. $o(g(n)) := \{f(n) : \text{for any } c > 0, \exists\, n_0 > 0 \mid 0 \le f(n) < c \cdot g(n) \text{ for all } n \ge n_0\}$

1.1.5 Little-ω

Definition 1.5. $\omega(g(n)) := \{f(n) : \text{for any } c > 0, \exists\, n_0 > 0 \mid 0 \le c \cdot g(n) < f(n) \text{ for all } n \ge n_0\}$

Theorem 1.2. $f(n) \in \omega g(n) \iff g(n) \in o(f(n))$

1.2 Proofs strategies

1.2.1 Proving from definition

This means that you'd use the definition to prove that a function g is in some complexity class.
Steps:

1.

Example:

1.2.2 Proving using limit properties

Use the following properties:

1.2.3 Helpful properties:

Transitivity:

$$
\begin{aligned}
f(n) = \Theta(g(n)) \text{ and } g(n) = \Theta(h(n)) &\implies f(n) = \Theta(h(n)) \\
f(n) = O(g(n)) \text{ and } g(n) = O(h(n)) &\implies f(n) = O(h(n)) \\
f(n) = \Omega(g(n)) \text{ and } g(n) = \Omega(h(n)) &\implies f(n) = \Omega(h(n)) \\
f(n) = o(g(n))) \text{ and } g(n) = o(h(n)) &\implies f(n) = o(h(n)) \\
f(n) = \omega(g(n)) \text{ and } g(n) = \omega(h(n)) &\implies f(n) = \omega(h(n))
\end{aligned}
$$

Reflexivity:

$$
\begin{aligned}
f(n) &= \Theta(f(n)) \\
f(n) &= O(f(n)) \\
f(n) &= \Omega(f(n))
\end{aligned}
$$

Symetry:

$$
f(n) = \Theta(g(n)) \iff g(n = \Theta(f(n)))
$$

Transpose symetry:

$$
\begin{aligned}
f(n) = O(g(n)) &\iff g(n) = \Omega(f(n)) \\
f(n) = o(g(n)) &\iff g(n) = \omega(f(n))
\end{aligned}
$$

Limit laws:

$$
\begin{aligned}
\lim_{n \to \infty} f(n)/g(n) \in [0, \infty) &\Rightarrow f = O(g) \\
\lim_{n \to \infty} f(n)/g(n) \in (0, \infty] &\Rightarrow f = \Omega(g) \\
\lim_{n \to \infty} f(n)/g(n) \in (0, \infty) &\Rightarrow f = \Theta(g) \\
\lim_{n \to \infty} f(n)/g(n) = 0 &\Rightarrow f = o(g) \\
\lim_{n \to \infty} f(n)/g(n) = \infty &\Rightarrow f = \omega(g)
\end{aligned}
$$

More:

$$
f_1 = O(g_1) \text{ and } f_2 = O(g_2) \implies f_1 + f_2 = O(\max(g_1, g_2))
$$

1.2.4 Recurrences (substitution, recurrence trees, master theorem)

1.3 Master method

Apply Master theorem to solve recurrances of the form $T(n) = aT(n/b) + f(n)$ where $a \geq 1$, $b > 1$, and that $f(n)$ is asymptotically nonnegative.

Master Theorem: fix this later

Theorem 1.3. Case 1: If $f(n) = O(n^{\log_b a - \epsilon})$ for some constant $a > 0$, then $T(n) = n \log^b a$.

Intuition: When $f(n)$ is polynomially less than the watershed function, the cost of leaves dominates the cost of the rest of the tree.

Case 2: If $f(n) = O(n \log^b a \log^k n)$ for some constant $k > 0$, then $T(n) = n \log^b a \log^{k+1} n$.

Intuition: There are $\log n$ levels, each with a cost of $n \log^b a \log kn$.

Simple (common) case: $k = 0 \Rightarrow f(n) = n \log^b a \Rightarrow T(n) = n \log^b a \log n$.

Case 3: If $f(n) = n \log^b a+$ for some constant $a > 0$, and $f(n)$ satisfies the regularity condition (below), then $T(n) = f(n)$.

Regularity condition: $af(n/b)^c \leq f(n)$ for some constant $c < 1$ and all large enough n. (Usually true for polynomial $f(n)$.)

Intuition: When $f(n)$ is polynomially greater than the watershed $f(n)$, the cost of the root dominates the cost of the leaves.

1.4 Recurrence trees

If of the form $T(n) = aT(n/b) + f(n)$, the each node has a branches and the starting a branches will be n/b. Then next layer will be n/b^2 because you are plugging n/b for n in $T(n)$.

The first layer is $f(n)$ time to combine, next layer will be $a/$ the denominator in that branch all times $f(n)$, basically the watershed function when you plug the new n in

keep going till you get 1. depth $= \log n$ base whatever n was divided by. make it into a sum

1. given recurrance relation, prove that it's in some class

2. method: (recurrence tree) -¿ guess, -¿ substitution method

3. look on recurrence sheet

4. ref sheet will have mast theorem, properties or relations for exponents and logarithms, summations

1.4.1 Models of computation (comparison model, random access model)

1.5 Solving problems by reducing to searching (choose + apply a data structure)

2 Interfaces & implementations

Interfaces: set, sequence, priority queue
Implementations: arrays, sorted arrays, linked lists, hash tables, BSTs, red-black trees, min/max heap

2.0.1 Priority queue interface

1. heap

2.1 Solving problems by reducing to sorting (choose + apply a sorting algorithm)

1. Mergesort Insertion sort Selection sort Heapsort

2.2 Solving problems by designing a new algorithm

1. Brute force "Decrease and conquer" Divide and conquer

3   Introduction

3.1   Vocabulary

1. interface:

2. implementation

3. array-based (example: array/list)

4. static: region of memory stays static

5. pointer-based (ex: linked list)

6. dynamic

4   Data Structures

sequence data structute operations in O

Implementation	build(X)	get_at/set_at	insert_first/delete_first	insert/delete last	insert at/delete at
array	n	1	n	n	n
linked list	n	n	1	n	n
dynamic array	n	1	n	$1_{(a)}$	n
binary tree	$n \log n$	h	h	h	h

Set data structute operations in O

Implementation	build(X)	find(k)	insert(x)/delete(x)	find_min/max()	find_prev/next
array	n	n	n	n	n
Sorted array	$n \log n$	n	n	1	$\log n$
direct access array	u	1	1	u	u
Hash table	$n_{(e)}$	$1_{(e)}$	$1_{(a)(e)}$	n	n
binary tree	n	h	h	h	h

4.1   Sequences

Given a collection of elements called X, store elements $x_0, \ldots, x_{n-1}$

4.1.1   Static Operations

1. build(X) - builds the sequence from X

2. iter_seq() output items in order

3. len() - returns the number of elements in X

4. get_at(i) - returns item at index i

5. set_at(i, x) - insert x at location i

4.1.2 Dynamic Operations

Dynamic sequences use static sequence operations in addition to these operations:

1. `insert_at(i, x)` - inserts item x at location i

2. `delete_at(i)` - deletes and returns item at location i

4.2 Sets

4.2.1 Static Operations

1. `build(X)` - builds set

   - with array implementation, `build(X)` $\in O(n)$ where $n = |X|$

2. `len()` - returns $|X|$

3. `find(k)` - returns item with key k

   - with array implementation, `find(k)` $\in O(n)$ where $n = |X|$

4. `find_min()` - returns item with smallest key

5. `find_max()` - returns item with largest key

6. `find_next(k)` - return next smallest item from k (?)

7. `find_prev(k)` - return next largest item from k (?)

8. with array, min, max, prev, next all $\in O(n)$

4.2.2 Dynamic Operations

1. `insert_at(x, i)` - insert element x at index i

2. `delete_at(i)` - delete and return element at i

3. with array, insert/delete $\in O(n)$

4.2.3 Running times

1. with sorted array, `find(k)`, `prev`, `next` $\in O(\log n)$, insert/delete in n, find min/max in 1, build is nlogn

4.3 Direct access table

4.3.1 Example

1. items are Reed students

2. "keyspace" is all possible 8-digit ID numbers

3. index by Reed ID makes `find(k)` $\in O(1)$, but wastes space

4.4   Hash tables

Before: sorting and searching using only comparisons on items.

Now: use keys in more complex ways arrow direct access table

```
h : {ids} -> {0, ... , m-1}
# assume m = Theta(m)
```

store items as determined by h(key)

cannot be a bijection

collision: things in ids that map to same idem in codomain

solution: instead of storing item in table, store pointer to another data structure

"Chaining" to some other data structure

```
def hash.find(k):
digest = h(k)
c = get(h.direst) -> go find k?
```

4.4.1   Running time

ideally: all chains are constant size so then all operations on chains are constant size

chose m,h such that all chains end up with about n/m approx O(1) things in them

4.4.2   Example

modular division: `h(k) = k mod(m)`.

Problem: only as even as keys themselves

if too even, might get mapped to same thing

4.4.3   Better example: universal family of hash functions

example of one family:

defien $h_{ab}(k) = (((ak + b) \mod p) \mod m)$ where p is a large prime and a, b chosen randomly from 0 to p-1

Formally, define family $\mathcal{H}(p, m) = \{h_{ab}(k) \mid a, b \in [0 \ldots p - 1]$ and $a \neq 0\}$ desired property: for some $h \in \mathcal{H}$,

Pr that any pairs of keys collide $[h(k_1) = h(k_2)] \leq \frac{1}{m}$ for all k1, k2 such that k1 not equal k2

parallel concept for sorting comparison model

lower bound Omega (n log n)

4.5   Binary Trees

traversal order: flatten tree **Operations:**

1. `find_first()`: finds the first node in traversal order. go left until you can't anymore

2. `find_last()`: finds last node in traversal order. goes right until you can't anymore

5 Running Time

5.1 red black

root black
every leaf black
if node is red, both children are black
for every node, all paths from node to each descendant have same number of black nodes

6 Correctness proof

6.1 Loop invariant method

Sketch:

1. Define a loop invariant L

2. **Initialization:** prove L holds before 1st iteration

3. **Maintenance:** if L holds before loop, still true before next one

4. **Termination:** show that loop terminates and that L still holds

7 Recursion

running time for recursive algo "solving a recurrance" - substitution method - recurrence tree - master theorem

8 Limits

8.1 Limit Laws

9 Sorting

Sorting Algorithm	Best Case	Average Case	Worst Case
Bubble Sort	$O(n)$	$O(n^2)$	$O(n^2)$
Selection Sort	$O(n^2)$	$O(n^2)$	$O(n^2)$
Insertion Sort	$O(n)$	$O(n^2)$	$O(n^2)$
Merge Sort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$
Quick Sort	$O(n \log n)$	$O(n \log n)$	$O(n^2)$
Heap Sort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$

9.1 Permutation Sort

```
permutation_sort(A):
    generate every possible permutation of A
    check if A is sorted
    continue until you find a sorted permutation
```

9.2 Selection Sort

```
selection_sort(A, i):
    1. finds biggest item in A[:i]
    2. swaps biggest thing with thing at A[i]
    3. recurse on A[:i-1]
```

9.3 Merge Sort

```
if n = 0, 1: done
otherwise:
    split array into Left, Right
    MergeSort(Left), MergeSort(Right)
    Merge(Left, Right)
```

1. specificity: array A, p starting point, r ending point

2. if p geq r (if p equal to r, then only one element and return)

3. convention: A[a:b] = [] if b less than A

4. a:b means including both end points

5. define midpoint $q = (p + r)/2$ (floor of the averge)

6. Left = A[p:q]

7. Right A[q+1:r]

8. MergeSort(A, p, q)

9. MergeSort(A, q+1, r)

10. Merge(A, p, q, r)

1. comparisons in constant time?

    (a) mostly comparing integers

    (b) assume $i \leq j$ in constant time

    (c) might be more complex

    (d)

10 Binary Search Trees

O(h) time where h is height
usually use these for sets
example: set us cs profs
name/id
adam/560

charlie/703

erica/998

greg/997

jim/100

include bst from notes

properties: for any X, keys in x.left leq x.key leq keys x.right

can also store a sequence in this way

ex 2: grocery list

apples, bananas, cereal, dish soap, eggs

order doesn't matter

bst property: for any L[i], all items in L[i].left appear before L[i]

all items in right subtree appear after L[i] is balanced, height is at most O(log n)

when balanced, called red-black trees

10.1   non-modifying operations

1. find min - all the down left

2. find max - all the way down right

3. kind k - compare at every node

4.

10.2   modifying properties

1. insert item: traverse, then insert as a leaf

2. delete: find successor and swap 16 with 17, then remove old leaf (find a safe leaf to remove)

3.

balanced if even under dynamic operations it maintains height O log N

red-black tree

1. dummy nodes

2. imagine single dummy node "tree ends"

3. stop drawing them

loop invariants

different cases for red black trees

11 priority queues and heaps

11.1 Priority Queue Interface

new: delete max - remove item with highest priority

find max = find highest priority item

supports subset of set interface operations

optemized for finding min or max, we will focus on max priority

also has insert and build

11.2 priority queue sort

algo in slides

simple because you make the ds do it

11.3 binary heaps

11.3.1 infer tree from aray

require: completely full in upper levels (until last level)

binary trees with n things -¿ arrays of length n

bottom level should be left justified

diferent than traversal ordering function

root is index 0

finding the left child of i is just at index 2i+1

right(i) = 2i+2

parent(i) = floor((i-1)/2)

max heap properties:

1. at node l, the value of thing at position l is going to be at least the value its children

2.

tree does not live in memory, just array

we want higher importance nodes to be closer to the top

bottom layer is less than all elements above it

example 1 from slides: not a max heap: 15 is a node that is is not greater than or equal to both of its children

inserting(x)

-append x to array: whole function is in O(logn)

-fix problems in heap

fixing: swap up with parents until max heap rule is followed:

check that val of parent geq val node, if no, swap and keep going

delete-max():

this deletes the root!

1. swap root with item at node number n-1, then delete thing at n-1

the fix downwards:

- start at root

- swap it with the greater key in its children, then continue

this is o of log n time

12 Cheat Sheets

12.1 Running Times for DS

Operations $O(\cdot)$				
Set Data Structure	Container	Static	Dynamic	Order
array	AX	ALA	248	
sorted array	AL	ALB	008	

12.2 orders

13 servers

1. servers and clients

2. each server: some capacity

3. each client: some workload

4. assign so that no server goes over its capacity

implementation

1. initialize(X) whre X is set of servers

2. connect(client, workload)

3. disconnect(client)

4. clients-of(server): return list of clients connected to that server

how to do:
idea
use priority queue for clients and prioritize biggest
priority queue for servers so that you can find best fit the best

notes on actual:
server has set of clients where key is client id and value is workload
priority queue storing the servers
store based on priorty, which is available capacity

structures:
pq P:(server.id, available capacity)
set S: (maps server id to set of clients and pointer to si in P) set C: maps client id to connected server and
workload (if not stored as client object)

14 Graph algorithms

14.1 Definition

1. $G = (V, E)$ where for all $(u, v) \in E$, $u \neq v$ and no duplicate edges (for this class)

2. $Adj + (u) := \{v \in V \mid (u, v) \in E\}$ (outgoing neighbors of node u)

3. incoming neighbord of u Adj - (u) = set(v in V st (v, u) in E)

4. out degree(u) = size of adj+

5. in degree(u) = size adj -

14.2 Representing graphs

matrix or adj list
storing:
top level set adj of vertices
lower level: adj lists Adj(u) either set or sequence (default to outgoing
can be direct access array
)

14.3 Running time

1. linear means O(V + E) (size of those sets)

   doing something with each edge and each vertex

2. $|E| = O(V^2)$ for simple graphs

   most edges you can have is fully connected graph

3. undirected: —E— leq size of V choose 2 (proved on homework)

4. directed : —E— leq 2 times (V size chose 2)

14.4 Breadth-first search

14.5 Path problems

15 Rules

15.1 Sum rules

15.2 Log rules

Product rule	$\log_b(MN) = \log_b(M) + \log_b(N)$
Quotient rule	$\log_b\left(\dfrac{M}{N}\right) = \log_b(M) - \log_b(N)$
Power rule	$\log_b(M^p) = p \cdot \log_b(M)$
Base switch rule	$\log_b(c) = \frac{1}{\log_c(b)}$
Base change rule	$\log_b(x) = \frac{\log_c(x)}{\log_c(b)}$

16   Midterm review

17   Lecture 16

17.1   Topological Sort

"finishing order" - order in which is finishes visiting (recursive calls pop all the way up)
DAG meaning?

> **Definition 17.1.** Topological order of $G = (V, E)$: ordering of vertices such that every (u, v) in E satisfies that u is before v in the order (in context of directed graphs)

Find topological ordering of graph: Edges = (A,C) (C, E) (A, B) (B, G) (B, E) (C, D) (E, D) (E, F)
check: ACBEDFG
other: ABGCEDF
order for F and G don't matter
Story problem: teams
Solution sketch:

1. Given info, what is the graph? vertices = set of employees, edges are links

2. goal: a topological ordering satisfies the constraint

3. find topological ordering or prove none exists

4. prove none exists: get sequence, check if topological

5. (reverse of visiting order is a topological sort)

6. solution exists iff G is a DAG

7. alternatively show a cycle so there is no ordering

8. or show graph is acyclic -¿ topological ordering exists

can solve by cycle detection problem

17.2   Weighted Graphs

unweighted: distance is num edges in shortest paths
add a weight function from set of edges to integers

17.3   weighted paths terminology

why infinum? has to do with range of weights we're trying to consider
    problem with negative weghts and cycles
confused...
    get nice graph

18 Lecture 17

assume G is (undirected) graph satisfies def 1 (show that G must also satisfy def 2)

weighted graphs given shortest paths weights, can compute shortest paths tree in linear time $(O(V + E))$

18.1 DAG Relaxation

18.1.1 high level:

for each vertex, maintain estimate which is upper bound as true shortest distance distance

repeatedly improve

refining estimates: true shortest-path weight between any pair of vertices u, v is upper bounded by weight of shortest path from u to v that has a specific last edge (x, v)

proposing path through x as possible path

see slide for formal definition

update estimate to use edge if edge exists to make it shorter

18.1.2 DAG Relaxation Algorithm

```
DAG-Relaxation_SP(G,s):
1. Initialize estimates:
   d(s,v) = inf for all v in V (excluding s)
   d(s,s) = 0
2. For each u in G in topological order:
   For each outgoing neighbor v in Adj+(u)

   ect
```

18.1.3 Correctness of DAG

Correctness condition:

Proof sketch: show .....

Base case: vertex s and every vertex before s in topological order satisfies claim at the start (because everything is assigned a distance of infinity)

inductivve step: assume claim holds

let pi be some shortest path from s to v, and let (u,v) be the last edge of pi

u must be beore v in top order, by induction, we already processed u, so by induction, estimate is true distance

after processing u:

1. see slides 2. see slides

input: graph G and s (any vertex, sometimes called source even though it's not a true source always) output: mapping from v in V to true distances from s, v

exercises:

Exercise 1: given undirected graph G, return whether G contains a negative-weight cycle

return true if there exists an edge such that w(edge) is negative

19 Lecture 18

19.1 Bellman-Ford

cycles and neg weights ok, so need to do neg weight cycle deal with

idea: more relaxation: while we can improve our estimate, keep improving

edge is tense if $d(s,v) > d(s,u) + w(u,v)$ for some u definitions: see slides

claim: after at most V-1 iterations, bellmanford halts with correct edge weights

subclaim: after i iterations, d(s,v) leq deltai()

19.2 notation

$\delta(u,v)$: "true" min weight from u $\rightarrow$ v

$d(u,v)$: current estimate of min weight from u $\rightarrow$ v

$w(u,v)$: weight of the edge from u $\rightarrow$ v

19.3 Dijkstra's Algorithm

like weighted bfs but priority queue

distance estimate for priority queue

"changeable" priority queue to update the keys

```
DIJKSTRA(G, w, s):
    INITIALIZE-SINGLE-SOURCE(G, s)
    S = {}
    Q = {}
    for each vertex u in G.V
        INSERT(Q, u)
    while Q not empty:
        u = EXTRACT-MIN(Q)
        S = S + {u}
        for each vertex v in
```

what is running time as function of:

insert: V
extract-min: V
decrease-key: E

fib heap (best for Dijkstra) which makes running time $O(|E| + |V|log|V|)$

sometimes insert can be replaced with build depending on implementation

dynamic array

binary heap

Correctness:

20   Midterm Review

Some things to know:

$$n^2 \notin o(2n^2)$$
$$\log_2(n^2) \in \Theta(\log_2 n)$$
$$(\log_2(n))^2 \notin \Theta(\log_2 n)$$
$$\log_{10}(n) \in \Theta(\log_2(n))$$
$$\sqrt{n} \in \omega(log_2(n))$$
$$\sqrt{n} \in \Theta(n)$$

21   Logical Equivalence

Contrapositive	$P \implies Q$	$\neg Q \implies \neg P$